

# **CS 553: CLOUD COMPUTING**

## **Tera Sort Benchmarking**

**By**

**Saptarshi Banerjee (A20376116)**

**Vaibhav Aggarwal (A20374988)**

**Abhishek Vijhani (A20377670)**

## **Shared Memory Sort**

We have written the code for the Shared Memory Tera Sort in Java. Tera Sort is a popular benchmark that measures the amount of time to sort one terabyte of randomly distributed data on a given computer system. We have performed sorting on 128GB and 1 TB datasets. Using the key's ASCII Value, we have performed sorting. We have divided the original dataset into the block size (size of file/ Number of Temporary files). After that we have created unsorted chunks for each file and applied Quick Sort on these chunks, which provides me with sorted chunks as output. This part of the program has threads applied to it to parallelize the process of sorting. Then with the help of k-way external merge sort, I have merged all the Sorted Chunks of earlier phase into a final sorted output file.

## **Hadoop Tera Sort**

We have written the code for Hadoop Sorting in Java. In this we have used Map-Reduce based implementation. It contains of two phases Sort and Shuffle Phase, which internally sorts the data without using any code. We have provided the mapper with input as key and values, and the Sort and Shuffle phase will categorize the data and sort them according to its Key. The reducer is used to give the output of the Sort and Shuffle Phase. We have tested the Hadoop code on 128 GB, 1 TB

## **Spark Tera Sort**

We have written the code for Spark using the PySpark Shell, which uses Python. It also uses the Map-Reduce based implementation. When we install Spark on it automatically installs various features like R-Studio, Scala, Python. We have implemented sorting using Python sortByKey(), functionality, which sorts the data based on Key value (10 Bytes). Then we obtained the chunks of the sorted data.

Performance is done on AWS instance (i3.large), using Ubuntu. The Specification for i3.large is

<b>Instance Type</b>	i3.large (Ubuntu)
<b>Ram</b>	15.25GB
<b>Number of Cores</b>	2(vCPU)
<b>Storage</b>	1 x 0.475 NVMe SSD
<b>Networking Performance</b>	Upto 10 Gigabit
<b>Region</b>	Us-east-1 (N. Virginia)
<b>EBS Volume Used</b>	500GB SSD

Performance is done on AWS instance (i3.4xlarge), using Ubuntu. The Specification for i3.4xlarge is

<b>Instance Type</b>	i3.4xlarge (Ubuntu)
<b>Ram</b>	122GB
<b>Number of Cores</b>	16(vCPU)
<b>Storage</b>	2 x 1.9 NVMe SSD
<b>Networking Performance</b>	Upto 10 Gigabit
<b>Region</b>	Us-east-1 (N. Virginia)
<b>EBS Volume Used</b>	2045GB SSD

# Shared Memory Tera Sort

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

Scheduled Instances

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

NETWORK & SECURITY

Security Groups

Launch Instance Connect Actions

Filter by tags and attributes or search by keyword

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)	IPv4 Public IP	IPv6 IPs
i-00374d033fe3ae028	i-00374d033fe3ae028	i3 large	us-east-1b	stopped	None	None	-	-	-
sap_test3	i-033852f33c56a2c24	i3 large	us-east-1b	stopped	None	None	-	-	-
sap_test4	i-05a363ad4244610...	i3 large	us-east-1b	running	2/2 checks ...	None	ec2-34-235-145-28.co...	34.235.145.28	-
vtest	i-063764b740a37a9f8	i3 xlarge	us-east-1b	running	2/2 checks ...	None	ec2-52-91-105-58.com...	52.91.105.58	-
sap_1tb	i-0d315c60a05c3391d	i3.4xlarge	us-east-1b	running	2/2 checks ...	None	ec2-54-84-87-111.comp...	54.84.87.111	-

Select an instance above

Feedback English (US) © 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

## i3.large: 128 GB Using gensort

```
root@ip-172-31-41-198: /vaibhav/64
File Edit View Search Terminal Help
-----
FILE PROCESSING STARTED...
-----
Please Wait ! This may take time.

1.Splitting Files into chunks
Free memory (bytes): 238655864
Unsorted Chunks created..
[You can find them in the Application's root directory as 'UNSORTED_CHUNKS']

Total no. of Chunks created:1822

Now Reading & Sorting Unsorted Chunks into Sorted Chunks. Please Wait...

Now Sorted Chunks will Sort and Merge into 1 File using Threads

Multi-Threading started...

Thread# 1 started..
Thread# 2 started..
Thread# 3 started..
Thread# 4 started..
Thread# 5 started..
Thread# 6 started..
Thread# 7 started..
Thread# 8 started..

Files Sorted and Merged successfully !!

Process Completed..
[Please find the file in the root folder as 'output.txt']

TOTAL TIME ELAPSED: 14981 seconds
```

Total time taken 14981 Seconds

### i3.4xlarge: 1 TB Using gensort

EBS volume used: 2045GB SSD

```
root@ip-172-31-32-52:/vaibhav/64
File Edit View Search Terminal Help
root@ip-172-31-32-52:/vaibhav/64# sudo rm -r SORTED_CHUNKS/
root@ip-172-31-32-52:/vaibhav/64# sudo rm -r output.txt
root@ip-172-31-32-52:/vaibhav/64# javac SharedMemory.java
root@ip-172-31-32-52:/vaibhav/64# java SharedMemory
-----
FILE PROCESSING STARTED...
-----
Please Wait ! This may take time.

1.Splitting Files into chunks
Free memory (bytes): 1909247144
Unsorted Chunks created..
[You can find them in the Application's root directory as 'UNSORTED_CHUNKS']

Total no. of Chunks created: 1035824

Now Reading & Sorting Unsorted Chunks into Sorted Chunks. Please Wait...

Now Sorted Chunks will Sort and Merge into 1 File using Threads

Multi-Threading started...

Thread# 1 started..
Thread# 2 started..
Thread# 3 started..
Thread# 4 started..
Thread# 5 started..
Thread# 6 started..
Thread# 7 started..
Thread# 8 started..

Files Sorted and Merged successfully !!

Process Completed..
[Please find the file in the root folder as 'output.txt']

TOTAL TIME ELAPSED: 65712 seconds
^Z
```

**Total time taken 65712 Seconds**

# Hadoop Tera Sort

## Configuring Hadoop on AWS

Here we have created i3.large instance on AWS

1. Login to new EC2 instance using ssh  
**(ssh -i aws-key.pem [ubuntu@54.183.30.236](mailto:ubuntu@54.183.30.236))**
2. Login as root user to install base packages (java 8)  
**(sudo su**  
**sudo apt-get install python-software-properties**  
**sudo add-apt-repository ppa:webupd8team/java**  
**sudo apt-get update**  
**sudo apt-get install oracle-java8-installer)**
3. Check the java version  
**(java –version)**
4. Download latest stable Hadoop using wget from one of the [Apache mirrors](#).  
**wget <http://www.trieuvan.com/apache/hadoop/common/hadoop-2.6.0/hadoop-2.6.0.tar.gz>**  
**(tar xzf hadoop-2.6.0.tar.gz)**
5. Create a directory where the hadoop will store its data. We will set this directory path in hdfs-site.  
**(mkdir hadoopdata)**
6. Add the Hadoop related environment variables in your bash file.  
**(vi ~/.bashrc**  
**export HADOOP\_HOME=/home/ubuntu/hadoop-2.6.0**  
**export HADOOP\_COMMON\_LIB\_NATIVE\_DIR=\$HADOOP\_HOME/lib/native**  
**export HADOOP\_OPTS="-Djava.library.path=\$HADOOP\_HOME/lib"**  
**export JAVA\_HOME=/usr/lib/jvm/java-8-oracle**  
**PATH=\$PATH:\$JAVA\_HOME/bin:\$HADOOP\_HOME/bin:\$HADOOP\_HOME/sbin")**
7. Save and exit and use this command to refresh the bash settings.  
**(source ~/.bashrc)**

8. Setting hadoop environment for password less ssh access. Password less SSH Configuration is a mandatory installation requirement. However it is more useful in distributed environment.

```
(ssh-keygen -t rsa -P ""  
cat $HOME/.ssh/id_rsa.pub >> $HOME/.ssh/authorized_keys)
```

9. check password less ssh access to localhost  
**(ssh localhost)**

10. exit from inner localhost shell  
**(exit)**

11. Set the hadoop config files. We need to set the below files in order for hadoop to function properly.

core-site.xml  
hadoop-env.sh  
yarn-site.xml  
hdfs-site.xml  
mapred-site.xml

12. go to directory where all the config files are present (cd /home/ubuntu/hadoop-2.6.0/etc/Hadoop), Copy and paste the below configurations in core-site.xml  
(Add the following text between the configuration tabs.)

```
<property>  
<name>hadoop.tmp.dir</name>  
<value>/home/ubuntu/hadooptmp/hadoop-${user.name}</value>  
<description>A base for other temporary directories.</description>  
</property>  
<property>  
<name>fs.default.name</name>  
<value>hdfs://localhost:9000</value>  
</property>)
```

13. Copy and paste the below configurations in hadoop-env.sh

get the java home directory using:

```
(readlink -f `which java` )
```

Example output: /usr/lib/jvm/java-8-oracle/jre/bin/java (NOTE THE JAVA\_HOME PATH.  
JUST GIVE THE BASE DIRECTORY PATH)

Need to set JAVA\_HOME in hadoop-env.sh  
**(export JAVA\_HOME=/usr/lib/jvm/java-8-oracle)**

14. Copy and paste the below configurations in mapred-site.xml

```
copy mapred-site.xml from mapred-site.xml.template  
(cp mapred-site.xml.template mapred-site.xml  
vi mapred-site.xml)
```

Add the following text between the configuration tabs.

```
<property>  
<name>mapred.job.tracker</name>  
<value>localhost:9001</value>  
</property>
```

15. Copy and paste the below configurations in yarn-site.xml

Add the following text between the configuration tabs.

```
<property>  
<name>yarn.nodemanager.aux-services</name>  
<value>mapreduce_shuffle</value>  
</property>
```

- Copy and paste the below configurations in hdfs-site.xml

Add the following text between the configuration tabs.

```
<property>  
<name>dfs.replication</name>  
<value>1</value>  
</property>  
<property><name>dfs.name.dir</name>  
<value>file:///home/ubuntu/hadoopdata/hdfs/namenode</value>  
</property>  
<property>  
<name>dfs.data.dir</name>  
<value>file:///home/ubuntu/hadoopdata/hdfs/datanode</value>  
</property>
```

16. Formatting the HDFS file system via NameNode (after installing hadoop, for the first time we have to format the HDFS file system to make it work)

**(hdfs namenode -format)**

17. Issue the following commands to start Hadoop

```
(cd sbin/  
./start-dfs.sh  
./start-yarn.sh)
```

**(start-all.sh)**

OR you can separately start required services as below:

**Name node:**

**(hadoop-daemon.sh start namenode)**

**Data node:**

**(hadoop-daemon.sh start datanode)**

**Resource Manager:**

**(yarn-daemon.sh start resourcemanager)**

**Node Manager:**

**(yarn-daemon.sh start nodemanager)**

**Job History Server:**

**(mr-jobhistory-daemon.sh start historyserver)**

18. Once hadoop has started point your browser to <http://localhost:50070/>

19. Check for hadoop processes /daemons running on hadoop with Java Virtual Machine Process Status Tool.

**(jps)**

20. OR you can check TCP and port details by using

**(sudo netstat -plten | grep java)**

# Configuring Hadoop on a Single Node

Here we are configuring the bashrc file

```
# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
#if [ -f /etc/bash_completion ] && ! shopt -oq posix; then
#    . /etc/bash_completion
#endif
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_INSTALL=/sapto/hadoop-2.7.4
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
```

Now we will configure xml file

```
<configuration>
<property>
<name>hadoop.tmp.dir</name>
<value>/home/ubuntu/hadooptmp/hadoop-${user.name}</value>
<description>A base for other temporary directories.</description>
</property>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

Configuring the HDFS file

Here we have kept the dfs replication 1

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property><name>dfs.name.dir</name>
<value>file:///sapto/hadoopdata/hdfs/namenode</value>
</property>
<property>
<name>dfs.data.dir</name>
<value>file:///sapto/hadoopdata/hdfs/datanode</value>
</property>
</configuration>
```

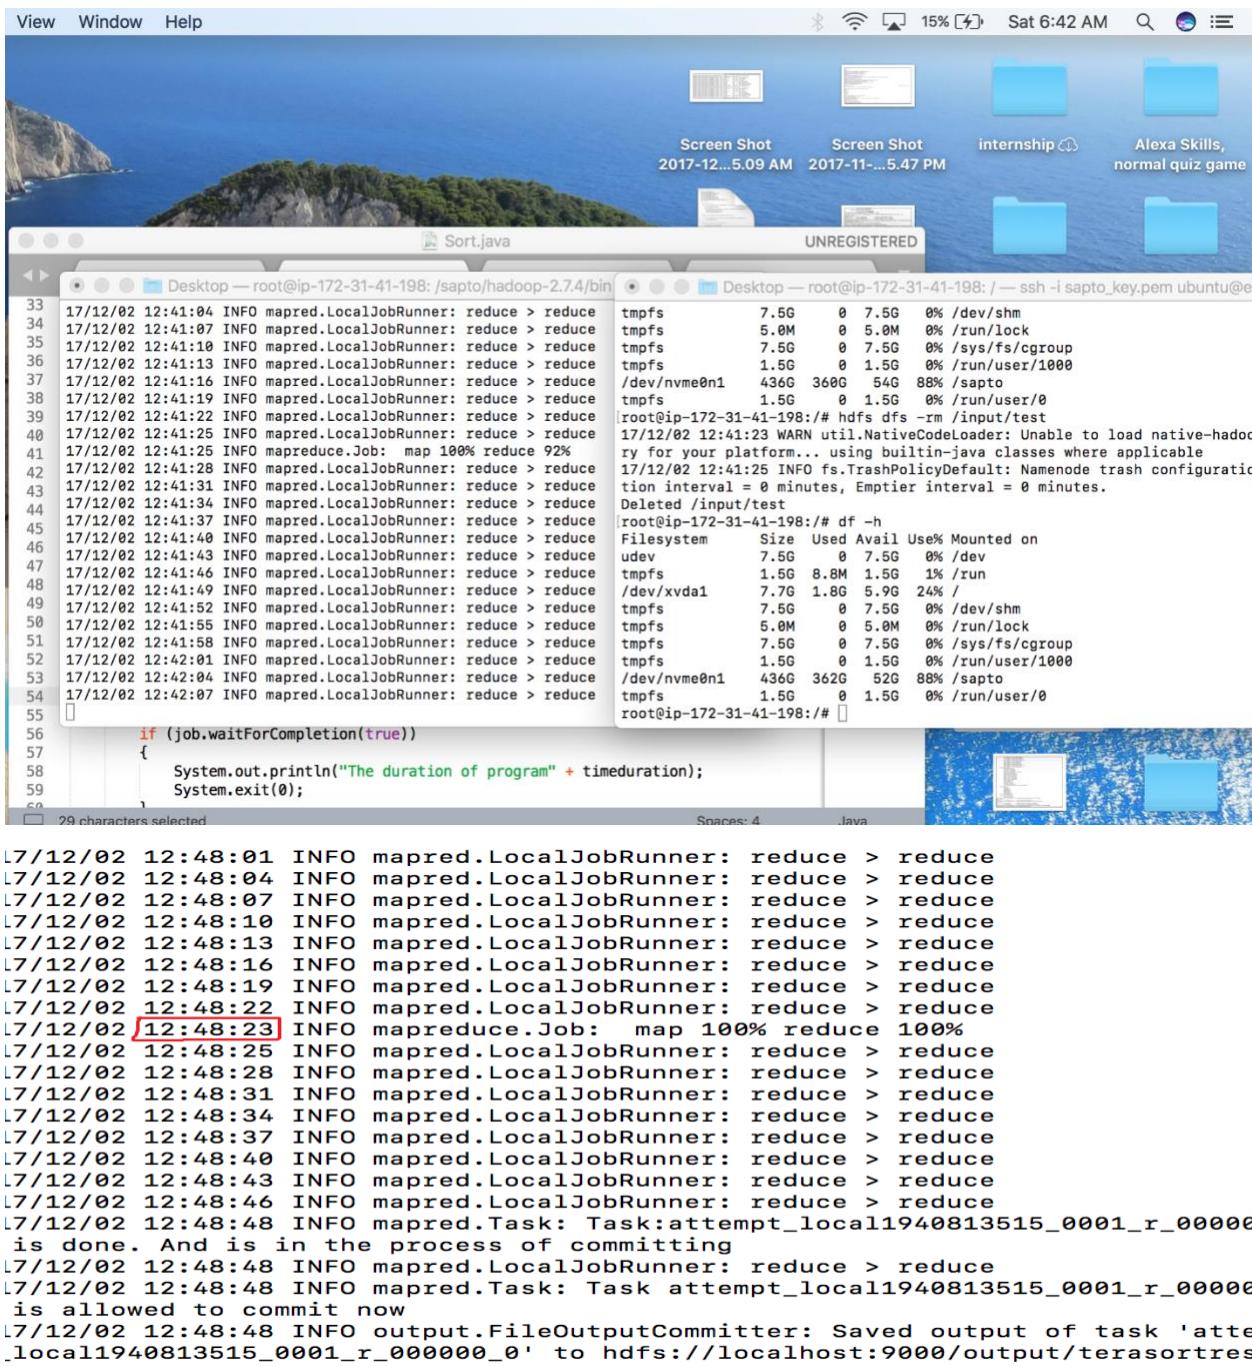
## Hadoop Single Node 128GB

**In this image, you can see the starting of Hadoop Job**

```
at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
root@ip-172-31-41-198:/sapto/hadoop-2.7.4/bin# hadoop jar sort1.jar SortHadoop
input/test3 /output/terasortresult
L7/12/02 10:23:11 [REDACTED] WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
L7/12/02 10:23:11 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id
L7/12/02 10:23:11 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=
L7/12/02 10:23:12 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
L7/12/02 10:23:12 INFO input.FileInputFormat: Total input paths to process : 1
L7/12/02 10:23:12 INFO mapreduce.JobSubmitter: number of splits:745
L7/12/02 10:23:12 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1940813515_0001
L7/12/02 10:23:12 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
L7/12/02 10:23:12 INFO mapreduce.Job: Running job: job_local1940813515_0001
L7/12/02 10:23:13 INFO mapred.LocalJobRunner: OutputCommitter set in config null
L7/12/02 10:23:13 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 1
```

## Hadoop Job running on way

```
Desktop — root@ip-172-31-41-198: /sapto/hadoop-2.7.4/bin [1] 17/12/02 12:04:11 INFO mapreduce.Job: map 100% reduce 38% 17/12/02 12:04:13 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:16 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:19 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:22 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:25 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:28 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:31 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:34 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:37 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:40 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:43 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:44 INFO mapreduce.Job: map 100% reduce 39% 17/12/02 12:04:46 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:49 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:52 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:55 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:58 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:01 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:04 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:07 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:10 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:13 INFO mapred.LocalJobRunner: reduce > sort Desktop — root@ip-172-31-41-198: / — ssh -i sapto_key.pem ubuntu@ 17/12/02 12:04:11 INFO mapreduce.Job: map 100% reduce 38% 17/12/02 12:04:13 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:16 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:19 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:22 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:25 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:28 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:31 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:34 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:37 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:40 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:43 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:44 INFO mapreduce.Job: map 100% reduce 39% 17/12/02 12:04:46 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:49 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:52 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:55 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:04:58 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:01 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:04 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:07 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:10 INFO mapred.LocalJobRunner: reduce > sort 17/12/02 12:05:13 INFO mapred.LocalJobRunner: reduce > sort tmpfs 1.5G 0 1.5G 0% /run/user/0 [root@ip-172-31-41-198: ~]# df -h Filesystem Size Used Avail Use% Mounted on udev 7.5G 0 7.5G 0% /dev tmpfs 1.5G 8.8M 1.5G 1% /run /dev/xvda1 7.7G 1.8G 5.9G 24% / tmpfs 7.5G 0 7.5G 0% /dev/shm tmpfs 5.0M 0 5.0M 0% /run/lock tmpfs 7.5G 0 7.5G 0% /sys/fs/cgroup tmpfs 1.5G 0 1.5G 0% /run/user/1000 /dev/nvme0n1 436G 304G 111G 74% /sapto tmpfs 1.5G 0 1.5G 0% /run/user/0 [root@ip-172-31-41-198: ~]# df -h Filesystem Size Used Avail Use% Mounted on udev 7.5G 0 7.5G 0% /dev tmpfs 1.5G 8.8M 1.5G 1% /run /dev/xvda1 7.7G 1.8G 5.9G 24% / tmpfs 7.5G 0 7.5G 0% /dev/shm tmpfs 5.0M 0 5.0M 0% /run/lock tmpfs 7.5G 0 7.5G 0% /sys/fs/cgroup tmpfs 1.5G 0 1.5G 0% /run/user/1000 /dev/nvme0n1 436G 306G 108G 74% /sapto tmpfs 1.5G 0 1.5G 0% /run/user/0 [root@ip-172-31-41-198: ~]
```



**Hadoop Job mapped and reduced 100%**

```

Combine input records=1000000000
Combine output records=1000000000
Reduce input groups=999999999
Reduce shuffle bytes=101000004470
Reduce input records=1000000000
Reduce output records=999999999
Spilled Records=4999604755
Shuffled Maps =745
Failed Shuffles=0
Merged Map outputs=745
GC time elapsed (ms)=106996
Total committed heap usage (bytes)=399716122624
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0

```

	size	Used	Avail	Use%	Mounted on
.5G	0	7.5G	0%	/dev/shm	
.0M	0	5.0M	0%	/run/lock	
.5G	0	7.5G	0%	/sys/fs/cgroup	
.5G	0	1.5G	0%	/run/user/1000	
36G	379G	36G	92%	/sapto	
.5G	0	1.5G	0%	/run/user/0	

```

-198:/sapto# df -h

```

	size	Used	Avail	Use%	Mounted on
.5G	0	7.5G	0%	/dev	
.5G	8.8M	1.5G	1%	/run	
.7G	1.8G	5.9G	24%	/	
.5G	0	7.5G	0%	/dev/shm	
.0M	0	5.0M	0%	/run/lock	
.5G	0	7.5G	0%	/sys/fs/cgroup	
.5G	0	1.5G	0%	/run/user/1000	
36G	332G	82G	81%	/sapto	
.5G	0	1.5G	0%	/run/user/0	

```

-198:/sapto# 

```

13515\_0001\_r\_000000\_0

17/12/02 12:48:48 INFO mapred.LocalJobRunner: reduce task executor complete.  
 17/12/02 12:48:50 INFO mapreduce.Job: Job job\_local1940813515\_0001 completed successfully

17/12/02 [REDACTED] 12:48:50 INFO mapreduce.Job: Counters: 35

#### File System Counters

```

FILE: Number of bytes read=38181308369819
FILE: Number of bytes written=75856220738849
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=37404037074944
HDFS: Number of bytes written=98999999901
HDFS: Number of read operations=559501
HDFS: Number of large read operations=0
HDFS: Number of write operations=748

```

#### Map-Reduce Framework

```

Map input records=1000000000
Map output records=1000000000
Map output bytes=99000000000
Map output materialized bytes=101000004470
Input split bytes=73010
Combine input records=1000000000
Combine output records=1000000000

```

**The Ending of Hadoop Job.**

## Screenshot of 128GB Output File

```

!0<v9`:`          0000000000000000000000000000000031290618 9999DDDDDEEE33331111AAAA2222
00005555AAAAEEEE00009999
!0<yzmJg          00000000000000000000000000000000275F95BF 0000AAAA44440000666677776666
0000BBBBB2222AAAA44440000
!0=#:U#q          00000000000000000000000000000001E724D94 FFFFAAAAEEEE5555EEEE33333333
EEEEAAAA77770000BBBBBDDDD
!0=)Z,x*          0000000000000000000000000000000027D70205 8888CCCC7777AAAA11118888EEE
FFFFF55556666777700002222
!0=06=XB          000000000000000000000000000000002915078B BBBBFFFF00004444000077770000
DDDD44443333BBBB99994444
!0=4nesb          000000000000000000000000000000002985F704 8888AAAA1111EEEE666600001111
BBBBB11114444EEEE8888DDDD
!0=<MM/ ]          00000000000000000000000000000001ED2AD1E AAAA8888FFFF8888CCCC3333BBBB
AAAAAFFFF33330000DDDD3333
!0=E1n(x          00000000000000000000000000000000267C9410 0000BBBB11111111111111AAAA8888
BBBBB999977779999BBBB1111
!0=IbC>|          000000000000000000000000000000003150E986 AAAADD5555222299998888EEE
CCCCBBBB3333EEEE0000BBBB
!0=Thpqa          00000000000000000000000000000000D26141D EEEE88882222444400006666EEE
8888FFFF5555CCCCAAAAAAA
!0=ei@b>          00000000000000000000000000000001A94D1F3 AAAA77773333000099999999FFF
AAAAAAAAEEEE99992222CCCC
!0=hK*t(          0000000000000000000000000000000022CC777A 333377779999DDDD777799998888
5555EEEEDDDDFFFFCCCC7777

```

<input type="checkbox"/>	i-00374d033fe3ae028	i3.large	us-east-1b	<span style="color: orange;">●</span> stopped	None		
<input type="checkbox"/>	sap_test3	i-033852f33c56a2c24	i3.large	<span style="color: orange;">●</span> stopped	None		
<input type="checkbox"/>	sap_test4	i-05a363ad4244610...	i3.large	<span style="color: green;">●</span> running	<span style="color: green;">✓</span> 2/2 checks ...	None	 ec2-34-235-145-28
<input type="checkbox"/>	vtest	i-063764b740a37a9f8	i3.xlarge	<span style="color: green;">●</span> running	<span style="color: green;">✓</span> 2/2 checks ...	None	 ec2-52-91-105-58.c
<input type="checkbox"/>	sap_1tb	i-0d315c60a05c3391d	i3.4xlarge	<span style="color: green;">●</span> running	<span style="color: green;">✓</span> 2/2 checks ...	None	 ec2-54-84-87-111.c

In this image you can see the total time elapsed on Hadoop Job

```
Combine output records=1000000000
Reduce input groups=100000000
Reduce shuffle bytes=10100000450
Reduce input records=100000000
Reduce output records=100000000
Spilled Records=395957882
Shuffled Maps =75
Failed Shuffles=0
Merged Map outputs=75
GC time elapsed (ms)=8255
Total committed heap usage (bytes)=40304115712
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=10000303104
File Output Format Counters
Bytes Written=99000000000
Total Elapsed Time on Hadoop: 972000000000
root@ip-172-31-41-198:/sapto/hadoop-2.7.4/bin#
```

So for 128 Gb it took us around 2.7 hrs.

## Hadoop Single Node 1TB

Here, you can see the starting time of Hadoop Job

```
root@ip-172-31-32-52:/sapto/hadoop-2.7.4/bin# hadoop jar sort.jar SortHadoop /in  
put/sortftb /output/sortftbout  
17/12/03 13:08:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
17/12/03 13:08:09 INFO Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session-id  
17/12/03 13:08:09 INFO jvm.JvmMetrics: Initializing JVM Metrics with processName=JobTracker, sessionId=  
17/12/03 13:08:09 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.  
17/12/03 13:08:09 INFO input.FileInputFormat: Total input paths to process : 1  
17/12/03 13:08:09 INFO mapreduce.JobSubmitter: number of splits:75  
17/12/03 13:08:09 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1478218440_0001  
17/12/03 13:08:09 INFO mapreduce.Job: The url to track the job: http://localhost:8080/  
17/12/03 13:08:09 INFO mapreduce.Job: Running job: job_local1478218440_0001  
17/12/03 13:08:09 INFO mapred.LocalJobRunner: OutputCommitter set in config null  
17/12/03 13:08:09 INFO output.FileOutputCommitter: File Output Committer Algorit
```

## Hadoop Job running for 1 TB File

```
Desktop — root@ip-172-31-32-52: /sapto/hadoop-2.7.4/bin — ssh -i sapto_ke...  
17/12/03 21:44:16 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:19 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:22 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:25 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:28 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:31 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:34 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:37 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:40 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:43 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:46 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:49 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:52 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:55 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:44:58 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:01 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:04 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:07 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:10 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:10 INFO mapreduce.Job: map 100% reduce 63%  
17/12/03 21:45:13 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:16 INFO mapred.LocalJobRunner: reduce > sort  
17/12/03 21:45:19 INFO mapred.LocalJobRunner: reduce > sort
```

### Here, we are generating sort file of 1TB

... omitted for brevity ...

u should have received a copy of the GNU General Public License  
ong with this program; if not, write to the Free Software Foundation,  
c., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301, USA.

```
rsion 1.5, cvs $Revision: 1.14 $  
MP Pump version 1.2.3, svn: 130  
ot@ip-172-31-32-52:/sapto1/64# ./gensort -a -t8 1000000000000 sortftb
```

### Here, Reduce Jobs are running for 1 Single node

```
17/12/03 22:44:53 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:44:56 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:44:59 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:02 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:05 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:08 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:11 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:14 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:17 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:20 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:23 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:26 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:29 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:32 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:35 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:38 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:41 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:44 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:47 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:50 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:53 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:56 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:45:59 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:02 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:05 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:08 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:11 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:14 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:17 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:20 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:23 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:26 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:29 INFO mapred.LocalJobRunner: reduce > reduce  
17/12/03 22:46:32 INFO mapred.LocalJobRunner: reduce > reduce
```

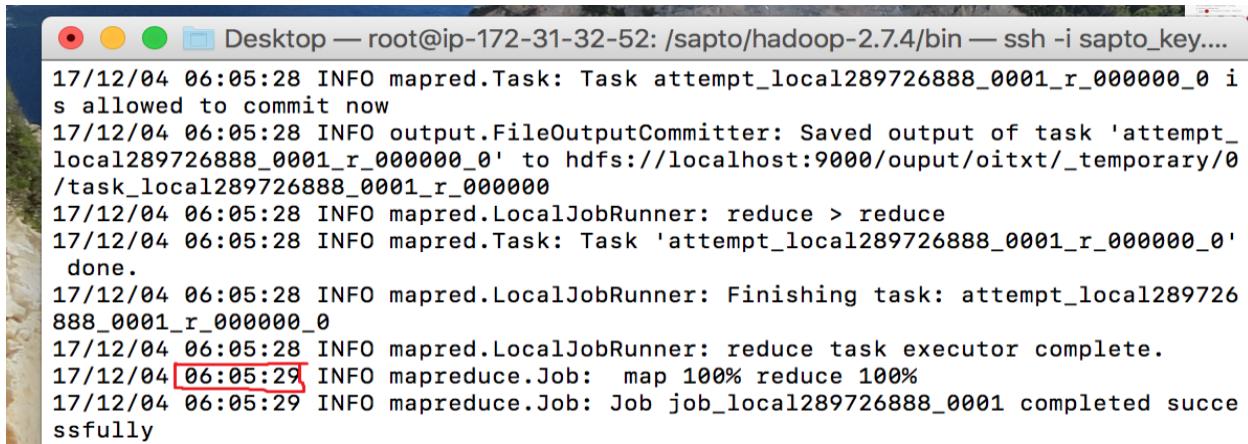
---

```

HDFS: Number of bytes written=296999999010
HDFS: Number of read operations=5013118
HDFS: Number of large read operations=0
HDFS: Number of write operations=2239
Map-Reduce Framework
  Map input records=3000000000
  Map output records=3000000000
  Map output bytes=297000000000
  Map output materialized bytes=303000013416
  Input split bytes=225836
  Combine input records=3000000000
  Combine output records=3000000000
  Reduce input groups=2999999990
  Reduce shuffle bytes=303000013416
  Reduce input records=3000000000
  Reduce output records=2999999990
  Spilled Records=16851737032
  Shuffled Maps =2236
  Failed Shuffles=0
  Merged Map outputs=2236
  GC time elapsed (ms)=383451
  Total committed heap usage (bytes)=1199164096512
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=300009154560
Total Time Elapsed time on Hadoop : 648000000000000
-----
```

**So, the total time required by Hadoop to sort 1 Tb data is around 18.5 hours**

Here, you can see the ending time of Hadoop Job



```

Desktop — root@ip-172-31-32-52: /sapto/hadoop-2.7.4/bin — ssh -i sapto_key.....
17/12/04 06:05:28 INFO mapred.Task: Task attempt_local289726888_0001_r_000000_0 is allowed to commit now
17/12/04 06:05:28 INFO output.FileOutputCommitter: Saved output of task 'attempt_local289726888_0001_r_000000_0' to hdfs://localhost:9000/ouput/oitxt/_temporary/0/task_local289726888_0001_r_000000
17/12/04 06:05:28 INFO mapred.LocalJobRunner: reduce > reduce
17/12/04 06:05:28 INFO mapred.Task: Task 'attempt_local289726888_0001_r_000000_0' done.
17/12/04 06:05:28 INFO mapred.LocalJobRunner: Finishing task: attempt_local289726888_0001_r_000000_0
17/12/04 06:05:28 INFO mapred.LocalJobRunner: reduce task executor complete.
17/12/04 06:05:29 INFO mapreduce.Job: map 100% reduce 100%
17/12/04 06:05:29 INFO mapreduce.Job: Job job_local289726888_0001 completed successfully
```

## Screenshot of 1TB Output File

## Configuring Hadoop on a Multi Node

Hadoop_slave3	i-005e39b3e38bde4...	r3.large	us-east-1c	<span>running</span>	<span>Initializing</span>
Hadoop_Master	i-02d844ff9ef6e75fd	i3.large	us-east-1a	<span>running</span>	<span>2/2 checks</span>
Hadoop_slave8	i-035bda34c050c8cdc	m3.xlarge	us-east-1d	<span>stopped</span>	
Hadoop_slave4	i-047488f90eb4c633c	r3.large	us-east-1c	<span>running</span>	<span>Initializing</span>
HadoopSlave1	i-04e202090f8de4fbc	i3.xlarge	us-east-1d	<span>running</span>	<span>Initializing</span>
	i-06336466af9221be	t2.micro	us-east-1c	<span>stopped</span>	
	i-066ed8a2ebaf45f8c	i3.large	us-east-1d	<span>stopped</span>	
Hadoop_slave7	i-08c347b92cb2dc447	i3.2xlarge	us-east-1d	<span>running</span>	<span>Initializing</span>
Hadoop_slave5	i-09895ff17b2853222	r3.large	us-east-1c	<span>running</span>	<span>Initializing</span>
	i-0998b1eacbb38b61	t2.medium	us-east-1c	<span>stopped</span>	
	i-0aee95c63eb9e8fb	t2.small	us-east-1c	<span>stopped</span>	
HadoopSlave2	i-0b987c5b78047e116	i3.4xlarge	us-east-1d	<span>running</span>	<span>Initializing</span>
sap_spark	i-0f827f3c045d8ffde	i3.large	us-east-1a	<span>stopped</span>	
Hadoop_slave6	i-0fa1e9dd0b7c18df5	r3.large	us-east-1c	<span>running</span>	<span>Initializing</span>

We can see Hadoop Master and Hadoop Slaves running. We configured the IPs of each and every slave.

Here we are configuring the .bashrc file

```
# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
#if [ -f /etc/bash_completion ] && ! shopt -oq posix; then
#    . /etc/bash_completion
#fi
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_INSTALL=/sapto/hadoop-2.7.4
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
```

Now we will configure xml file

```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://master:9001</value>
</property>
</configuration>
```

## Configuring the HDFS file

Here we have kept the dfs replication 1

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.permissions</name>
<value>false</value>
</property>
</configuration>
```

## Configuring the YARN file

Here, we have changed the port number of slaves

```
<?xml version="1.0"?>
<configuration>
  <!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>master:8038</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>master:8039</value>
  </property>
  <property>
    <name>yarn.resourcemanager.webapp.address</name>
    <value>master:8086</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>master:8037</value>
  </property>
  <property>
    <name>yarn.resourcemanager.admin.address</name>
    <value>master:8033</value>
  </property>
</configuration>
```

## Hadoop Multi Node 1TB

Here, we are generating the 1TB Data  
.....

You should have received a copy of the GNU General Public License along with this program; if not, write to the Free Software Foundation, c., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301, USA.

```
rsion 1.5, cvs $Revision: 1.14 $
MP Pump version 1.2.3, svn: 130
ot@ip-172-31-32-52:/sapto1/64# ./gensort -a -t8 1000000000000 sortfb
```

Here, you can see the ending time of the Hadoop Job

---

```
HDFS: Number of bytes written=296999999010
HDFS: Number of read operations=5013118
HDFS: Number of large read operations=0
HDFS: Number of write operations=2239
Map-Reduce Framework
    Map input records=3000000000
    Map output records=3000000000
    Map output bytes=297000000000
    Map output materialized bytes=303000013416
    Input split bytes=225836
    Combine input records=3000000000
    Combine output records=3000000000
    Reduce input groups=2999999990
    Reduce shuffle bytes=303000013416
    Reduce input records=3000000000
    Reduce output records=2999999990
    Spilled Records=16851737032
    Shuffled Maps =2236
    Failed Shuffles=0
    Merged Map outputs=2236
    GC time elapsed (ms)=383451
    Total committed heap usage (bytes)=1199164096512
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=300009154560
Total Time Elapsed time on Hadoop : 54000000000000
```

-----

Here the time taken by Hadoop to complete sorting of 1 TB using multi nodes is around 16 hours.

## Screenshot of 1TB Output File

## Configuring Spark on AWS

1. In a web browser from your local development environment, visit the Apache Spark Download Page. We need to generate a download link which we can access from our EC2 instance.

## Download Apache Spark™

1. Choose a Spark release:
2. Choose a package type:
3. Download Spark: [spark-2.2.0-bin-hadoop2.7.tgz](#)
4. Verify this release using the [2.2.0 signatures and checksums](#) and [project release KEYS](#).

*Note: Starting version 2.0, Spark is built with Scala 2.11 by default. Scala 2.10 users should download the Spark source package and build with Scala 2.10 support.*

The Download Link in the 3<sup>rd</sup> bullet dynamically updates based on your choices.

2. Right Click on the download link and copy it into your clipboard so it can be pasted onto your EC2 instance. From your EC2 instance , type these commands

```
#Download Spark to the ec2-user's home directory
cd ~
wget https://d3kbcqa49mib13.cloudfront.net/spark-2.2.0-bin-hadoop2.7.tgz
# Unpack Spark in the /opt directory
sudo tar zxvf spark-2.2.0-bin-hadoop2.7.tgz -C /opt
# Create a symbolic link to make it easier to access
sudo ln -fs spark-2.2.0-bin-hadoop2.7 /opt/spark
```

3. To complete the installation, set the SPARK\_HOME environment variable so it takes effect when you login into EC2 instance.

```
# Insert these lines into your ~/.bash_profile:
export SPARK_HOME=/opt/spark
PATH=$PATH:$SPARK_HOME/bin
export PATH
# Then exit the text editor and return to the command line.
```

4. You need to reload the environment variables

```
# Reload environment variables
source ~/.bash_profile
```

```
# Confirm that spark-submit is now in the PATH.  
spark-submit --version  
# (Should display a version number)
```

5. Spark is not shy about presenting reams of informational output to the console as it runs. To make the output of your applications easier to spot, you can optionally set up Log4j configuration file and suppress some of the spark logging output. You should maintain logging at the INFO level during the Sparkour tutorials, as we will review some information such as Master URLs in the log files.

```
# Create a Log4J configuration file from the provided template.  
cp $SPARK_HOME/conf/log4j.properties.template  
$SPARK_HOME/conf/log4j.properties  
vi $SPARK_HOME/conf/log4j.properties  
# (on line 19 of the file, change the log level from INFO to ERROR)  
# (Note that this will suppress some of the output needed in the tutorials)  
# log4j.rootCategory=ERROR, console  
# Save the file and exit the text editor.
```

# Configuring Spark on a Single Node

Here we are configuring the .bashrc file

```
# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
#if [ -f /etc/bash_completion ] && ! shopt -oq posix; then
#    . /etc/bash_completion
#fi
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_INSTALL=/sapto/hadoop-2.7.4
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export SPARK_HOME=/sapto/spark-2.2.0-bin-hadoop2.7
export PATH=$PATH:$SPARK_HOME/bin
```

Now we will configure xml file

```
<configuration>
<property>
<name>hadoop.tmp.dir</name>
<value>/home/ubuntu/hadooptmp/hadoop-${user.name}</value>
<description>A base for other temporary directories.</description>
</property>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

Configuring the HDFS file

Here we have kept the dfs replication 1

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property><name>dfs.name.dir</name>
<value>file:///sapto/hadoopdata/hdfs/namenode</value>
</property>
<property>
<name>dfs.data.dir</name>
<value>file:///sapto/hadoopdata/hdfs/datanode</value>
</property>
</configuration>
```

## Spark Single Node 128GB

Here, we are starting the Spark Job, you can see the start time of Spark in the image

```
to adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel
(newLevel).
17/12/03 [22:19:51] WARN NativeCodeLoader: Unable to load native-hadoop library for
your platform... using builtin-java classes where applicable
17/12/03 22:19:52 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting
port 4041.
Path of input file ->/sapto/64/sortfile
Path of output file ->/sapto/64/output
Stage 0:=====
(409 + 1) / 2981]
```

### Spark Output file for 128GB

```
*****
Value of B is "v*NX#W+/" 000000000000000000000000000000001D7B1 44446666EEEE8888999
9BBBB6666999988887777888811116666
*****
Value of B is "v4<73S,u 00000000000000000000000000000000434898 777788880000EEEE222
2AAA8888DDDEEEBBBBB111188889999
*****
Value of B is "v8%kmcD 00000000000000000000000000000000475597 9999FFFF44447777999
90000444422229999DDDDAAAAEEE8888
*****
Value of B is "v:BDC9aP 000000000000000000000000000000003BA1DE 666677779999AAA888
86666EEEECCCC6666888833311113333
*****
Value of B is "v:e<jPY~ 000000000000000000000000000000007C5502 222200001111BBBB666
633330000EEEE3333CCCC3333DDDFFFF
*****
Value of B is "v;hAL&t- 0000000000000000000000000000000036C993 BBBBFFFF2222CCCCFFF
F77773333999977770000BBBB8888CCCC
*****
Value of B is "v>_c+d}U 000000000000000000000000000000008817B0 5555000088889999444
4777766669999DDDD5555FFFF99991111
*****
Value of B is "vC.Dg<OJ 0000000000000000000000000000000033DD7F 8888222277778888CCC
CDDDDCCCCCCCC8888EEEE222288880000
```

**In this image, you can see the ending time of Spark**

**Total time taken for spark to sort 128 GB data is 1.2 Hours**

## Spark Single Node 1TB

Here, we are starting the Spark Job, you can see the start time of Spark in the image.

```
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
17/12/04 01:15:07 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
17/12/04 01:15:09 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
Path of input file ->/sapto1/64/sortftb
Path of output file ->/sapto1/output.txt
[Stage 0:>
[Stage 0:>                                (0 + 0) / 8941
[Stage 0:>                                (0 + 1) / 8941
[Stage 0:>                                (1 + 1) / 8941
[Stage 0:>                                (2 + 1) / 8941
[Stage 0:>                                (3 + 1) / 8941
[Stage 0:>                                (4 + 1) / 8941
[Stage 0:>                                (5 + 1) / 8941
]
[Stage 0:>                                (109 + 1) / 8941
[Stage 0:>                                (110 + 1) / 8941
[Stage 0:>                                (111 + 1) / 8941
[Stage 0:>                                (112 + 1) / 8941
[Stage 0:>                                (113 + 1) / 8941
[Stage 0:>                                (114 + 1) / 8941
[Stage 0:>                                (115 + 1) / 8941
[Stage 0:>                                (116 + 1) / 8941
[Stage 0:>                                (117 + 1) / 8941
[Stage 0:>                                (118 + 1) / 8941
[Stage 0:>                                (119 + 1) / 8941
[Stage 0:>                                (120 + 1) / 8941
[Stage 0:>                                (121 + 1) / 8941
[Stage 0:>                                (122 + 1) / 8941
[Stage 0:>                                (123 + 1) / 8941
[Stage 0:>                                (124 + 1) / 8941
[Stage 0:>                                (125 + 1) / 8941
[Stage 0:>                                (126 + 1) / 8941
[Stage 0:>                                (127 + 1) / 8941
[Stage 0:>                                (128 + 1) / 8941
[Stage 0:>                                (129 + 1) / 8941
[Stage 0:>                                (130 + 1) / 8941
[Stage 0:>                                (131 + 1) / 8941
[Stage 0:>                                (132 + 1) / 8941
[Stage 0:>                                (133 + 1) / 8941
[Stage 0:>                                (134 + 1) / 8941
[Stage 0:>                                (135 + 1) / 8941
[Stage 0:>                                (136 + 1) / 8941
[Stage 0:>                                (137 + 1) / 8941
[Stage 0:>                                (138 + 1) / 8941
[Stage 0:>                                (139 + 1) / 8941
[Stage 0:>                                (140 + 1) / 8941
[Stage 0:>                                (141 + 1) / 8941
[Stage 0:>                                (142 + 1) / 8941
```

In this image, you can see the ending time of Spark, and Output file of Spark

```
Desktop — ubuntu@ip-172-31-32-52: /sapto1/spark-2.2.0-bin-hadoop2.7/pyth...
Value of B is ~`I^1|=.$ 00000000000000000000000000000000FE59 CCCC00001111BBBB0000
4444111166664444AAAA88882222EEEE
*****
Value of B is ~`Wz(;<ij) 00000000000000000000000000000000E349 EEEE5555CCCC00000000
9999EEEE2222DDDD22228888BBBBEEEE
*****
Value of B is ~`]nb\{/Fc 00000000000000000000000000001008 8888DDDD77778888EEEE
88882222BBBBCCCC8888444455559999
*****
Value of B is ~`o',`V}G1 000000000000000000000000000000008AC7 00001111AAAA1111BBBB
CCCC44448888FFFF77771111BBBB8888
*****
Value of B is ~`q++>he m 000000000000000000000000000000005B0A BBBBDDDD88881111CCCC
BBBB1111DDDD3333FFFFFF77777777
*****
Value of B is ~`yuN\$@D 00000000000000000000000000000000E3C4 1111EEEE0000AAAA7777
88883333AAAA00001111AAACCCDDDD
17/12/04 09:02:10 ERROR Executor: Exception in task 0.0 in stage 0.0 (TID 0)
org.apache.spark.api.python.PythonException: Traceback (most recent call last):
  File "pyspark/worker.py", line 177, in main
```

**Total time taken is by spark to sort 1TB data is 9 Hours**

# Configuring Spark on a Multi Node

**Here we are configuring the bashrc file**

```
# enable programmable completion features (you don't need to enable
# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
#if [ -f /etc/bash_completion ] && ! shopt -oq posix; then
#    . /etc/bash_completion
#endif
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_INSTALL=/sapto/hadoop-2.7.4
export HADOOP_CLASSPATH=${JAVA_HOME}/lib/tools.jar
export PATH=$PATH:$HADOOP_INSTALL/bin
export PATH=$PATH:$HADOOP_INSTALL/sbin
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_HOME=$HADOOP_INSTALL
export HADOOP_HDFS_HOME=$HADOOP_INSTALL
export YARN_HOME=$HADOOP_INSTALL
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native
export SPARK_HOME=/sapto/spark-2.2.0-bin-hadoop2.7
export PATH=$PATH:$SPARK_HOME/bin
```

**Now we will configure xml file**

```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://master:9001</value>
</property>
</configuration>
```

## Configuring the HDFS file

Here we have kept the dfs replication 1

```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.permissions</name>
<value>false</value>
</property>
</configuration>
```

## Configuring the YARN file

---

```
<?xml version="1.0"?>
<configuration>
  <!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>master:8038</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>master:8039</value>
  </property>
  <property>
    <name>yarn.resourcemanager.webapp.address</name>
    <value>master:8086</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>master:8037</value>
  </property>
  <property>
    <name>yarn.resourcemanager.admin.address</name>
    <value>master:8033</value>
  </property>
</configuration>
```

## Spark Multi Node 1TB

Here, we are starting the spark job, you can see the starting time of spark in the image

```
root@ip-172-31-46-168:/sapto/spark-2.2.0-bin-hadoop2.7/python# python spark_sort.py /sapto/64/sortfile /sapto/output.txt
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
17/12/04 00:52:31 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
17/12/04 00:52:32 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
17/12/04 00:52:32 WARN Utils: Service 'SparkUI' could not bind on port 4041. Attempting port 4042.
Path of input file ->/sapto/64/sortfile
Path of output file ->/sapto/output.txt
[Stage 0:>                                         (5 + 1) / 2981]
```

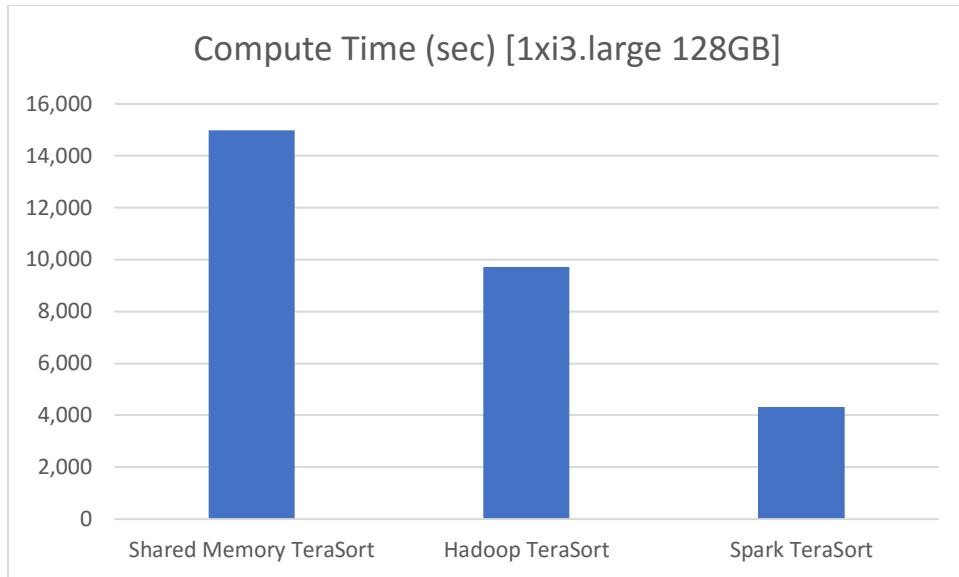
In this image you can see the ending time of spark.

```
Desktop — root@ip-172-31-46-168: /sapto/spark-2.2.0-bin-hadoop2.7/python...
Value of B is ~~%i?_ "?= 0000000000000000000000000000000013BF9 FFFF999977778888555
59999E44411114444666677778888EEEE
*****
Value of B is ~~FZ>[_0& 00000000000000000000000000000000490D AAAA9999FFFF0000DDD
DCCCCFFFF3333DDDD4444CCCC3333AAAA
*****
Value of B is ~~I^1]|=. 00000000000000000000000000000000FE59 CCCC00001111BBBB000
04444111166664444AAAA88882222EEEE
*****
Value of B is ~~Wz(;<iJ) 00000000000000000000000000000000E349 EEEE5555CCCC0000000
09999EEEE2222DDDD22228888BBBBEEEE
*****
Value of B is ~~]nb\{/Fc 0000000000000000000000000000000010008 8888DDDD77778888EEE
E88882222BBBBCCCC888844455559999
*****
Value of B is ~~o', 'V}G1 000000000000000000000000000000008AC7 00001111AAAA1111BBB
BCCCC44448888FFFF77771111BBBB8888
*****
Value of B is ~~q+*>he m 000000000000000000000000000000005B0A BBBBDDDD88881111CCC
CBBBB1111DDDD3333FFFFEEEE77777777
*****
Value of B is 17/12/04 07:24:59 ERROR Executor: Exception in task 0.0 in stage 0.
0 (TID 0)
org.apache.spark.api.python.PythonException: Traceback (most recent call last):
```

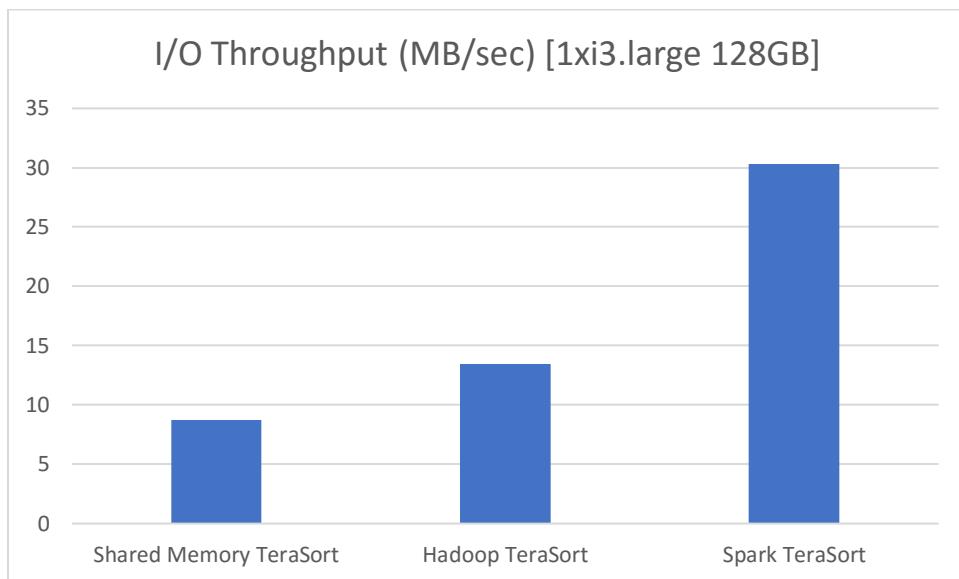
Total time taken by spark to sort 1 TB data in multimode is around 7 hours from the starting and ending time.

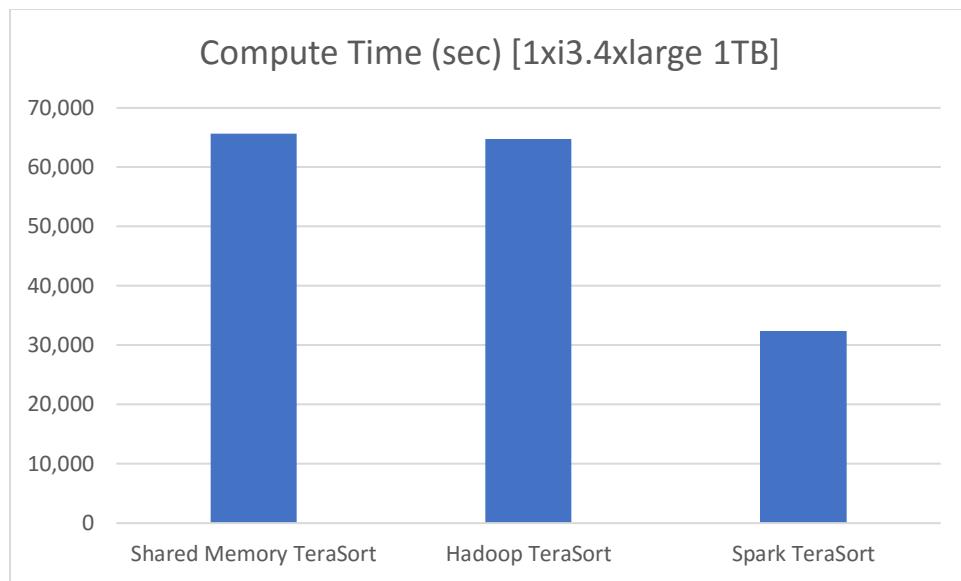
## PERFORMANCE EVALUATION

Experiment (instance/dataset)	Shared Memory Tera Sort	Hadoop Tera Sort	Spark Tera Sort
<b>Compute Time (sec) [1xi3.large 128GB]</b>	14,981	9720	4320
<b>Data Read (GB) [1xi3.large 128GB]</b>	128	128	128
<b>Data Write (GB) [1xi3.large 128GB]</b>	128	124	122
<b>I/O Throughput (MB/sec) [1xi3.large 128GB]</b>	8.74	13.48	30.3
<b>Compute Time (sec) [1xi3.4xlarge 1TB]</b>	65,712	64800	32400
<b>Data Read (GB) [1xi3.4xlarge 1TB]</b>	1,024	1,024	1,024
<b>Data Write (GB) [1xi3.4xlarge 1TB]</b>	1,024	988	975
<b>I/O Throughput (MB/sec) [1xi3.4xlarge 1TB]</b>	15.95	15.80	31.60
<b>Compute Time (sec) [8xi3.large 1TB]</b>	N/A	57,600	28,800
<b>Data Read (GB) [8xi3.large 1TB]</b>	N/A	1,024	1,024
<b>Data Write (GB) [8xi3.large 1TB]</b>	N/A	988	975
<b>I/O Throughput (MB/sec) [8xi3.large 1TB]</b>	N/A	17.77	35.5
<b>Speedup (Weak Scale)</b>	The speedup goes down as the number of thread increases. It is not affected for thread 2,3,4 but it significantly goes down for 7,8 threads.	2.2Mb/sec With increase from 1 node to 8 nodes, we observe that the speedup increases by 2.2 Mb/sec for Hadoop	5.2Mb/sec With increase from 1 node to 8 nodes, we observe that the speedup increases by 5.2Mb/sec
<b>Efficiency (Weak Scale)</b>	Efficiency is affected in same manner, down for more thread used.	Efficiency increases by 13.2%	Efficiency increases by 17.7%

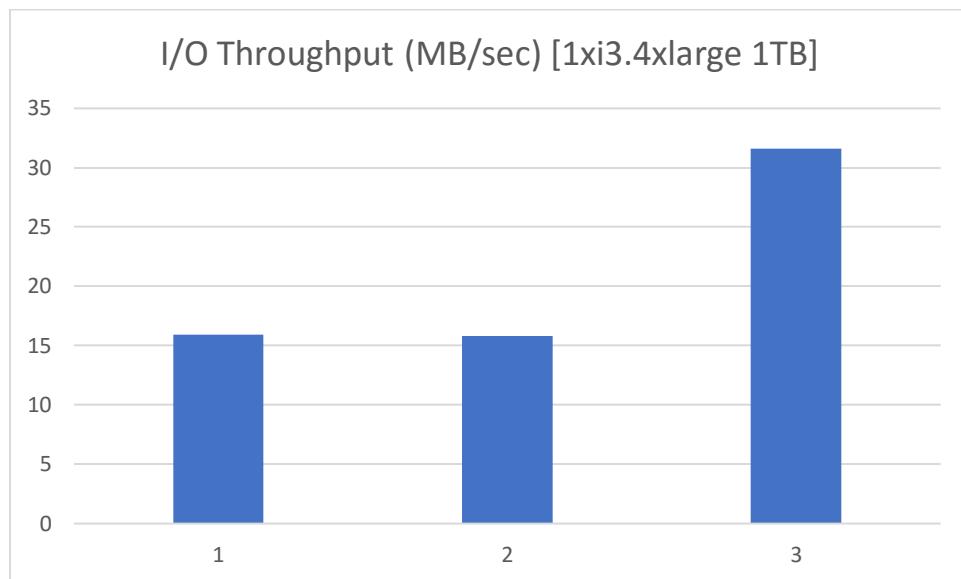


**The Graph clearly shows Spark is most efficient in sorting big volumes of data**

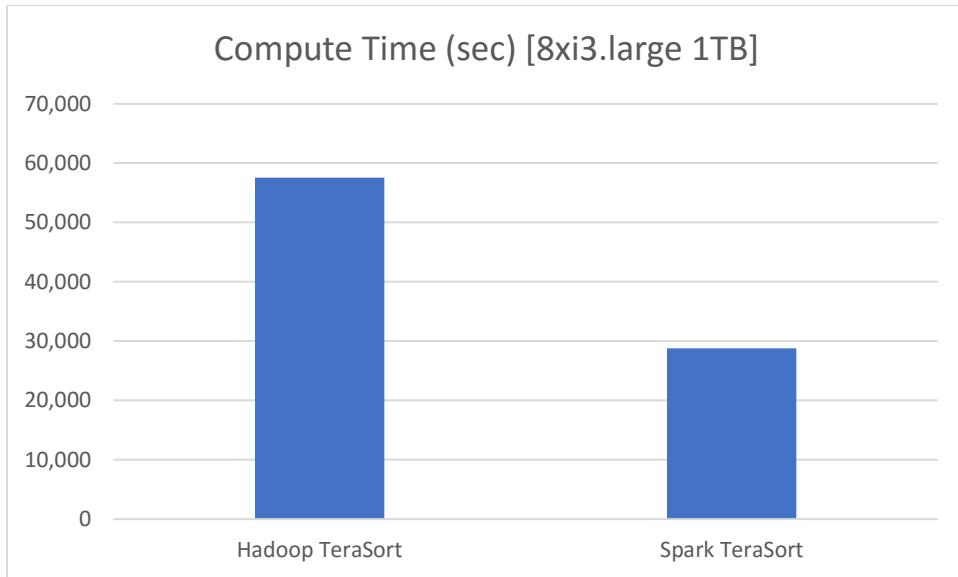




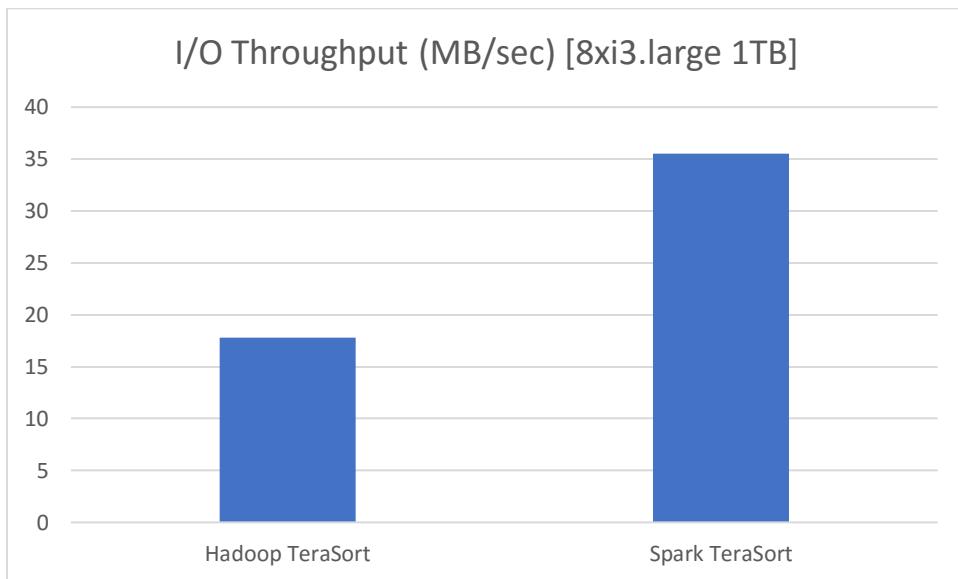
**Even for 1TB spark comes out to be the best, Hadoop second best and shared last.**



**Graph for Multi-node(Hadoop and Spark)**



**Also in the Multi-node 1TB sorting, spark is better**



## Conclusion

So, we draw the conclusion that Apache spark is the best of all for sorting big volumes of data. It is much faster and efficient than Hadoop. Hadoop is much better than Shared memory for sorting large volumes of data.

Spark seems to be the best for 1 node scale with a throughput of 30.3 Mb/Sec compared 13.48 Mb/sec and 8.74 Mb/sec for Hadoop and shared respectively.

In 8 Nodes, we see that spark is better and achieves a slightly better efficiency than Hadoop. According to me for 100 node scales or 1000 node scale if we use both apache Hadoop and spark together we will get the best results.

Our benchmarking closely aligns with the result with the winners using apache Hadoop and Spark. It is clearly evident apache spark is the winner. It is 100 times faster than Hadoop and very efficient.

All the big companies are now backing to use apache spark to perform big data operations.

## References

- <http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/>
- <https://www.edureka.co/blog/install-hadoop-single-node-hadoop-cluster>
- <http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-multi-node-cluster/>
- <http://grepcode.com/file/repo1.maven.org/maven2/org.apache.hadoop/hadoop-mapreduce-examples/2.6.0/org/apache/hadoop/examples/terasort/TeraGen.java>
- [https://github.com/hparik11/TeraSort\\_on\\_Hadoop-Spark/blob/master/Code/SortHadoop.java](https://github.com/hparik11/TeraSort_on_Hadoop-Spark/blob/master/Code/SortHadoop.java)
- <https://github.com/facebookarchive/hadoop-20/blob/master/src/examples/org/apache/hadoop/examples/terasort/TeraGen.java>
- <https://github.com/kunaldhande/Projects/blob/master/CS553%20-%20TeraSort/SourceCode/hadoop/HadoopSort.java>
- <https://hadoop.apache.org/docs/r2.7.1/api/org/apache/hadoop/examples/terasort/packager-summary.html>
- <http://www.chinacloud.cn/upload/2014-01/14010410467139.pdf>
- <http://www.dreamincode.net/forums/topic/223614-how-to-create-a-shell-script-to-compile-and-execute-a-java-file/>
- <https://stackoverflow.com/questions/14211076/reading-shared-memory-data-using-java-that-is-written-by-c>
- <https://github.com/shalinc/Cloud-Sorting-Large-Dataset/blob/master/Code/SharedMemoryScript.sh>
- <http://docs.aws.amazon.com/sdk-for-java/v1/developer-guide/setup-install.html>
- <http://www.ordinal.com/gensort.html>